

Deep learning

1. Introduction

A course @EDHEC
by Romain Tavenard (Prof. @Univ. Rennes 2)

Course details

- 18 hours (3 days)
- Instructor: Romain Tavenard romain.tavenard@univ-rennes2.fr
- Tools (*cf* course page for required Python packages):
 - Deepnote (create an account)
 - or Google Colab (create an account)
 - or Jupyter Notebooks running on your machine
- Evaluation
 - Group project (50%)
 - Final exam: A few questions + Lab session (50%)


Pre-requisites

- Basics of Python coding
- A (tiny) bit of calculus
 - What's the derivative of a function?
 - Functions of several variables
- Machine Learning topics
 - Empirical risk optimization and its limitations
 - Model evaluation & selection (cross-validation)

Our first model: the Perceptron

- Input: $x = \{x_0, \dots, x_D\}$
- Output: a
- Parameters to be optimized: $\{w_0, \dots, w_D, b\}$
- Activation function (chosen a priori): φ


$$a = \varphi \left(\underbrace{\sum_j w_j x_j + b}_o \right)$$



Motivating example

- Dataset: Boston housing prices
- Toy task: predict housing price (PRICE) based on average number of rooms per dwelling (RM)
- Chosen model: linear regression without intercept
$$\text{PREDICTED PRICE} = w_0 \times \text{RM}$$
- Cost function (also called loss): Mean Squared Error



$$\sum_i \underbrace{(\text{PREDICTED PRICE}_{(i)} - \text{PRICE}_{(i)})^2}_{w_0 \times \text{RM}_{(i)}}$$



Motivating example: Optimization

- Goal: Find w_0 that minimizes our loss
- Testing many values at random is not tractable
- Alternative strategy: Gradient descent

$$w^{(t+1)} \leftarrow w^{(t)} - \rho \nabla_w \mathcal{L}(w^{(t)})$$



Course outline

- Model architectures
 - Multi-Layer Perceptrons (for tabular data)
 - Convolutional models (for images)
 - Recurrent models & attention-based models (for time series)
- Different losses for different learning tasks
- Optimization strategies (variants of gradient descent)