

# Deep Learning

## 5. Time series & sequences


---

A course @EDHEC  
by Romain Tavenard (Prof. @Univ. Rennes 2)


# Neural network architectures and inductive biases

---


1D Conv.



RNN




Self-Attention



# Convolutional neural nets for time series

---




Source: [Le Guennec *et al.*, 2014]

# Recurrent neural nets

---


- Very flexible model (any length, let the model learn its memory needs, ...)



Source: [Christopher Olah's blog](#)

# Recurrent neural nets

- "Vanilla" RNN in more details



$$h_t = \tanh(W_h h_{t-1} + W_x x_t + b)$$





Illustration: RNN cell, source: [Christopher Olah's blog](#)

# Recurrent neural nets

---


- Very flexible model (any length, let the model learn its memory needs, ...)
- Difficult to learn in practice
  - Slow (lack of parallelism)
  - Vanishing gradients (hard to learn long-term dependencies)



Source: [Christopher Olah's blog](#)

# Recurrent neural nets

- Gated Recurrent Unit (GRU)
- Principle
  - At each time step, keep only part of the information




$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Illustration: GRU cell, source: [Christopher Olah's blog](#)



# Recurrent neural nets

---

- Long Short-Term Memory (LSTM)
- Principle: similar to GRU




Illustration: LSTM cell, source: [Christopher Olah's blog](#)

# Recurrent neural nets

---

- Long Short-Term Memory (LSTM)
- Principle: similar to GRU





Illustration: LSTM cell, source: [Christopher Olah's blog](#)

# Recurrent neural nets

---

- Long Short-Term Memory (LSTM)
- Principle: similar to GRU



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Illustration: LSTM cell, source: [Christopher Olah's blog](#)

# Recurrent neural nets

---

- Long Short-Term Memory (LSTM)
- Principle: similar to GRU





Illustration: LSTM cell, source: [Christopher Olah's blog](#)

# Recurrent neural nets

---

- Long Short-Term Memory (LSTM)
- Principle: similar to GRU



$$o_t = \sigma (W_o [ h_{t-1}, x_t ] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$

Illustration: LSTM cell, source: [Christopher Olah's blog](#)


# Attention-based models

## Motivating example

---

- Consider the following translation task:
  - From English: "**An apple** that had been on the tree in the garden for weeks had finally been **picked up**."
  - To French: "**Une pomme** qui était sur l'arbre du jardin depuis des semaines avait finalement été **ramassée**."
- Both recurrent and convolutional architectures fail at modelling long-range dependencies (for different reasons)
  - Attention aims at tackling this limitation

# Attention for Sequence-to-Sequence (seq2seq) tasks



## Recurrent models

- Compute a bottleneck representation
- Generate output sequence from this bottleneck

## Attention-based models

- For each token in the output sequence, aggregate input features
- Aggregation is importance-based
- Importance depends on the features, not their localization

# Assessing importance: Queries, Keys & Values


---

- The **Query / Key / Value** metaphor: Python dictionaries

```
d = {"a": 12, "b": 7}  
print(d["a"])
```
- Queries, Keys & Values: the continuous case
  - Look for keys that are **similar** to the query (not equal)
  - Retrieved value is a weighted average of values associated to keys that are close to the query
    - Could be seen as weighted k-nearest neighbors

# Visual explanation of the Attention mechanism


$$A_{i,j} = \text{softmax}(Q_i \cdot K_j)$$



# Visual explanation of the Attention mechanism


---

$$Y_i = \sum_j A_{i,j} V_j$$



# Visual explanation of the Attention mechanism


$$A_{i,j} = \text{softmax}(Q_i \cdot K_j)$$



# Visual explanation of the Attention mechanism

---


$$Y_i = \sum_j A_{i,j} V_j$$



# Standard attention layers

---


- Standard Attention layers
  - take as inputs 2 sequences  $X$  and  $X'$
  - output a sequence  $Y$



# Standard attention layers

---


- In self-attention layers, we have  $X=X'$




# Standard attention layers

---

- In multi-head attention layers,
  - several ( $h$  here) such blocks operate in parallel
  - Their output is concatenated in feature dimension



# The Transformer: The typical attention-based architecture



# Summary

---

- 1d-CNN, RNN and Attention-based models can be used
  - depends on the context
  - slightly different underlying assumptions
    - locality (ConvNets)
    - sequentiality (RNNs)
    - relationships between any set of items in the sequence (Attention-based models)
  - 1d-CNN are faster to train