

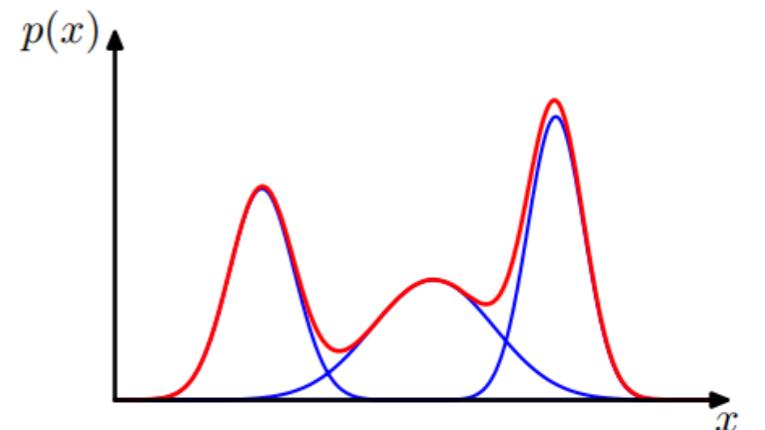
Deep Learning

6. Generative neural networks

A course @EDHEC
by Romain Tavenard (Prof. @Univ. Rennes 2)

Generative models in a nutshell

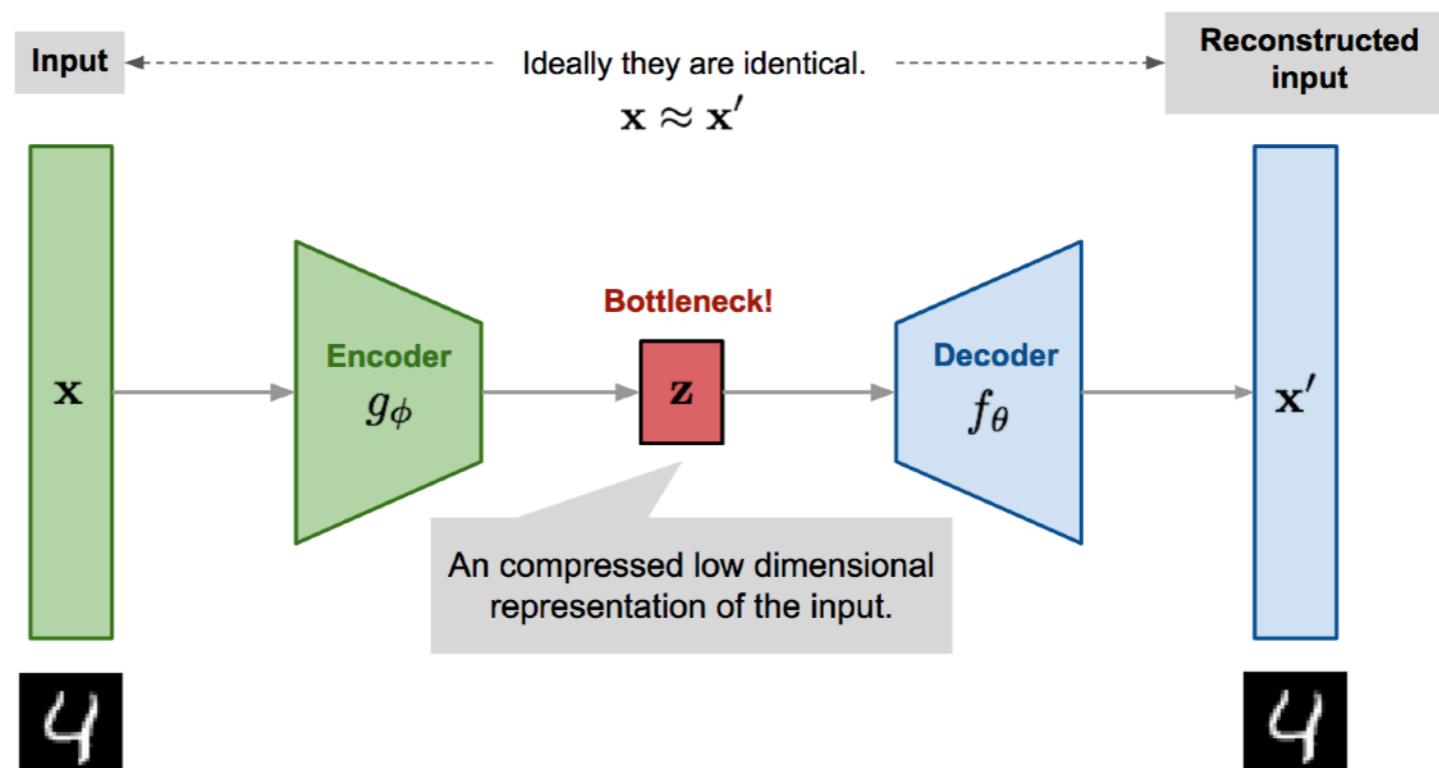
- Goal: model $p(x)$ or $p(x|y)$
 - explicitly
 - eg. Gaussian Mixture Models
 - **implicitly**
 - at least allow for sampling (*i.e.* generate new data)
 - e.g. Variational Auto Encoders,
Generative Adversarial Networks



Auto-encoders

[Hinton & Salakhutdinov, 2006]

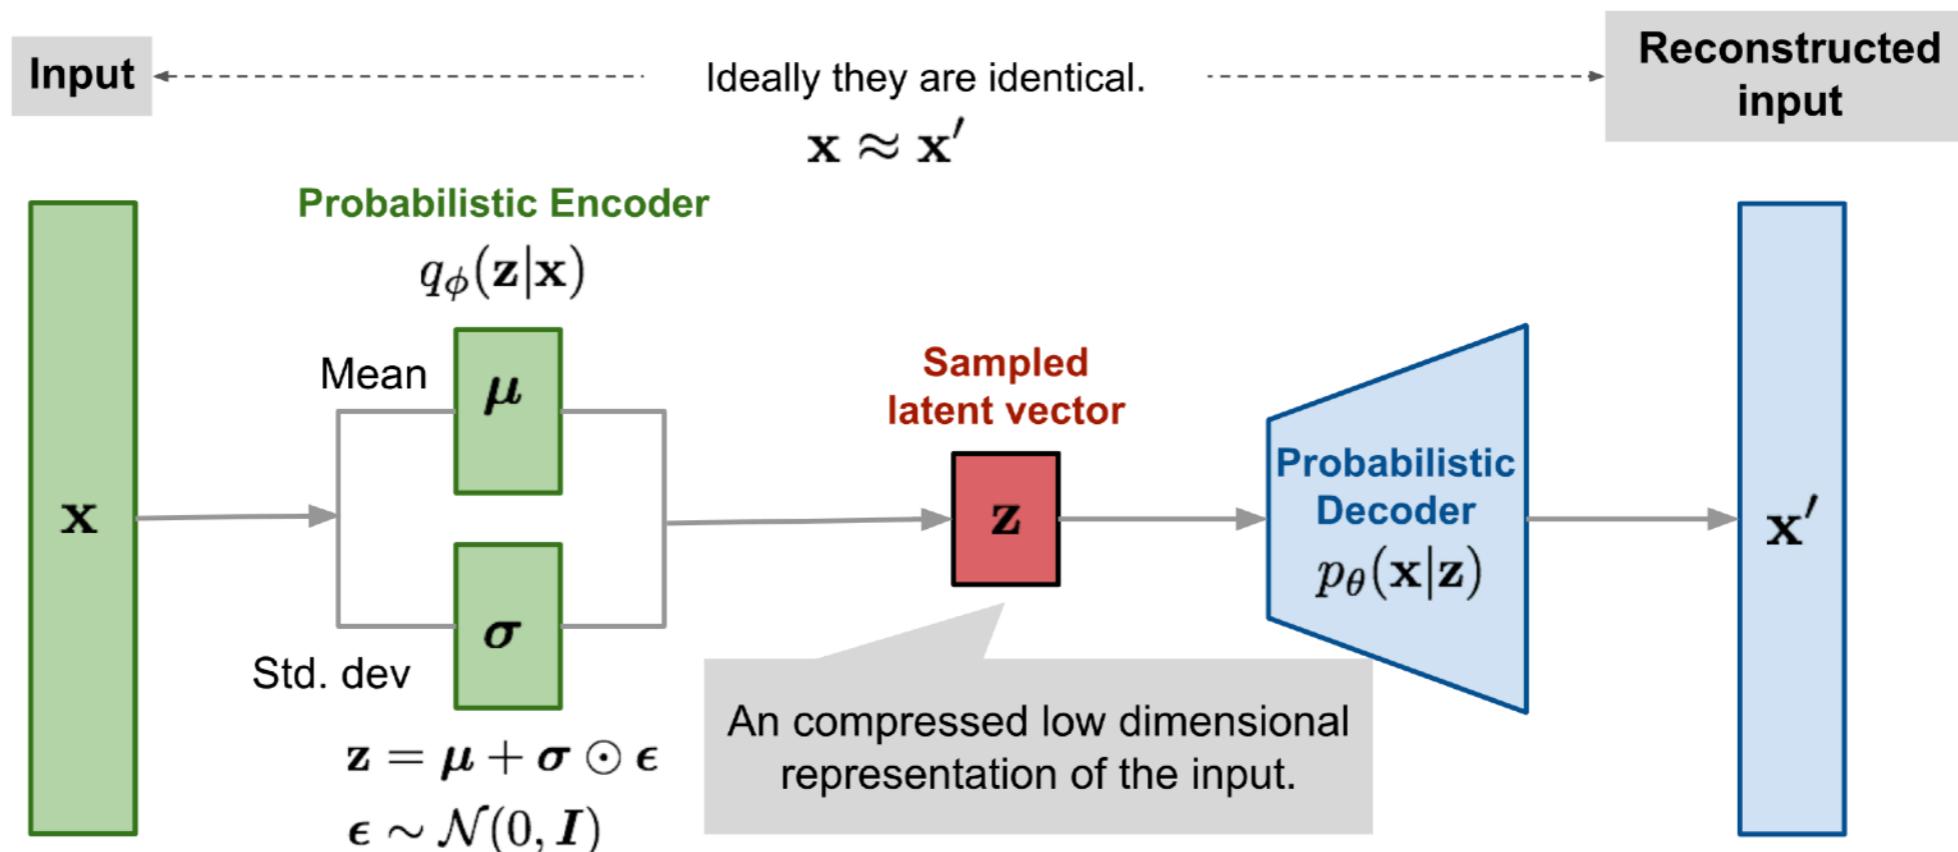
- Encode information in a latent space
 - typically lower-dimensional
 - similar in spirit to a non-linear PCA
 - not a generative model *per se!*



Variational auto-encoders

[Kingma & Welling, 2014]

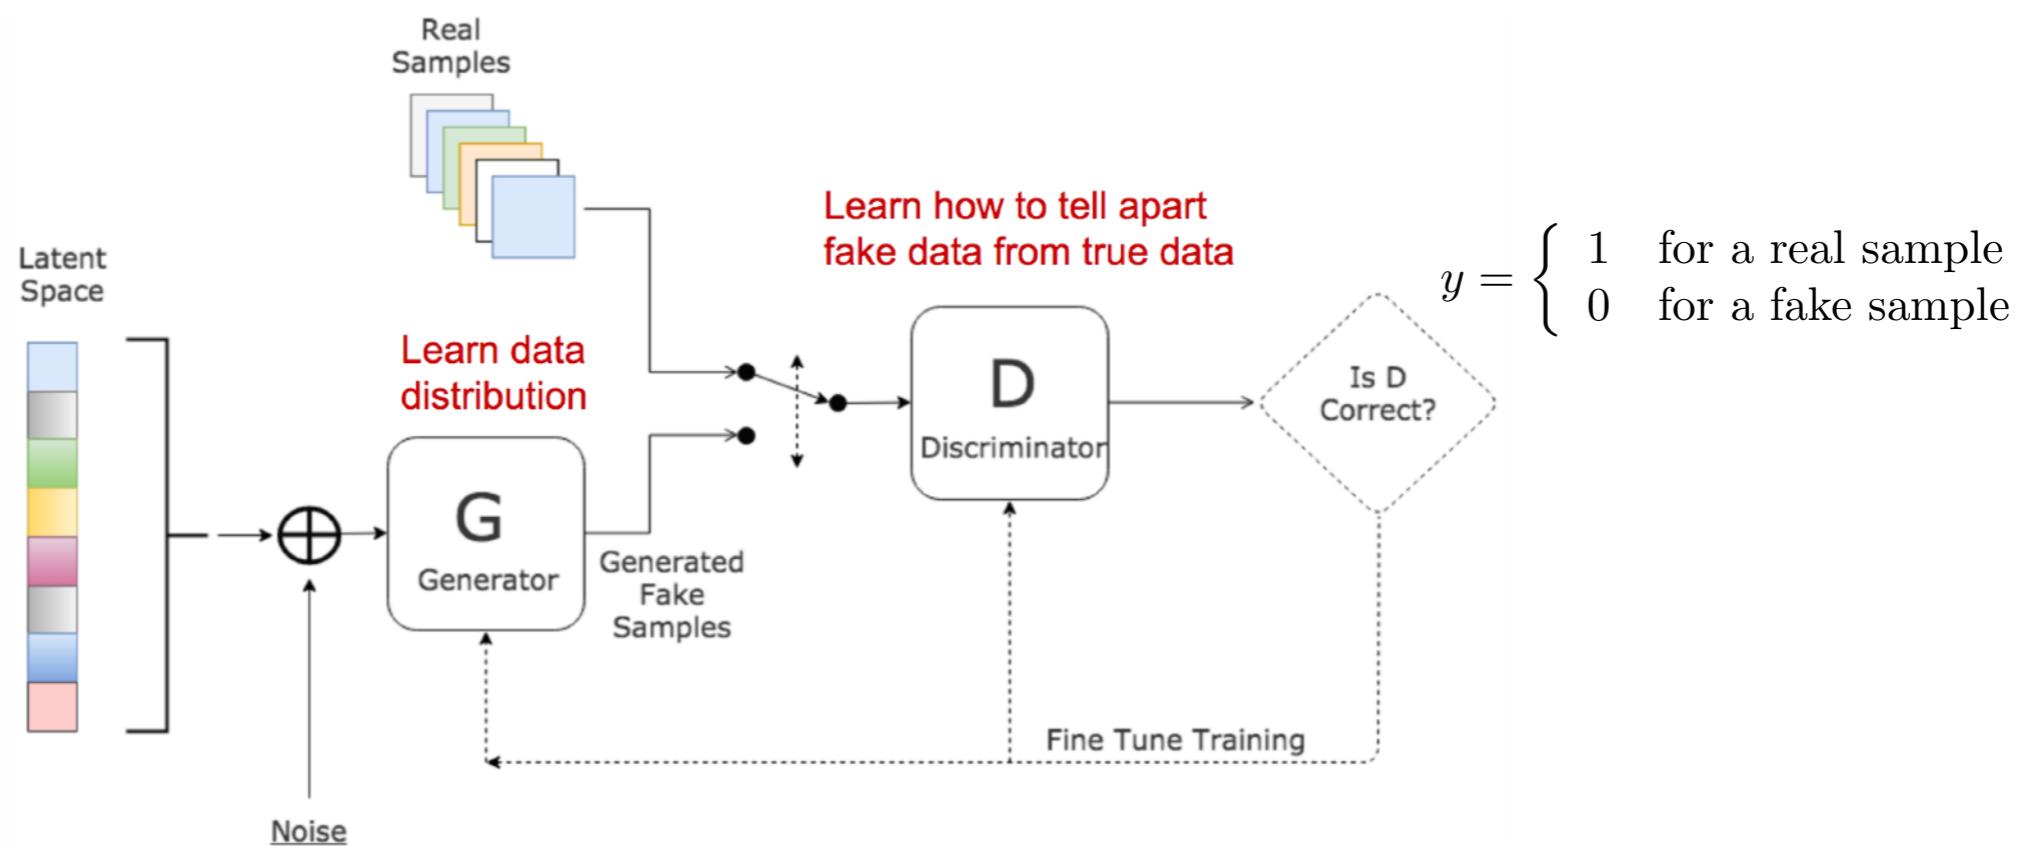
- Goal: turn auto-encoders into generative models
- Idea: set a prior on the distribution of z (through penalization of the loss function)
- Generative process: draw z from the prior, and decode it



Generative Adversarial Networks

[Goodfellow et al., 2014]

- Model:



Source: lilianweng.github.io

- Loss function:

$$\begin{aligned} \min_G \max_D L(D, G) &= \mathbb{E}_{x \sim p_r(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \\ &= \mathbb{E}_{x \sim p_r(x)} [\log D(x)] + \mathbb{E}_{x \sim p_g(x)} [\log(1 - D(x))] \end{aligned}$$

Generative Adversarial Networks

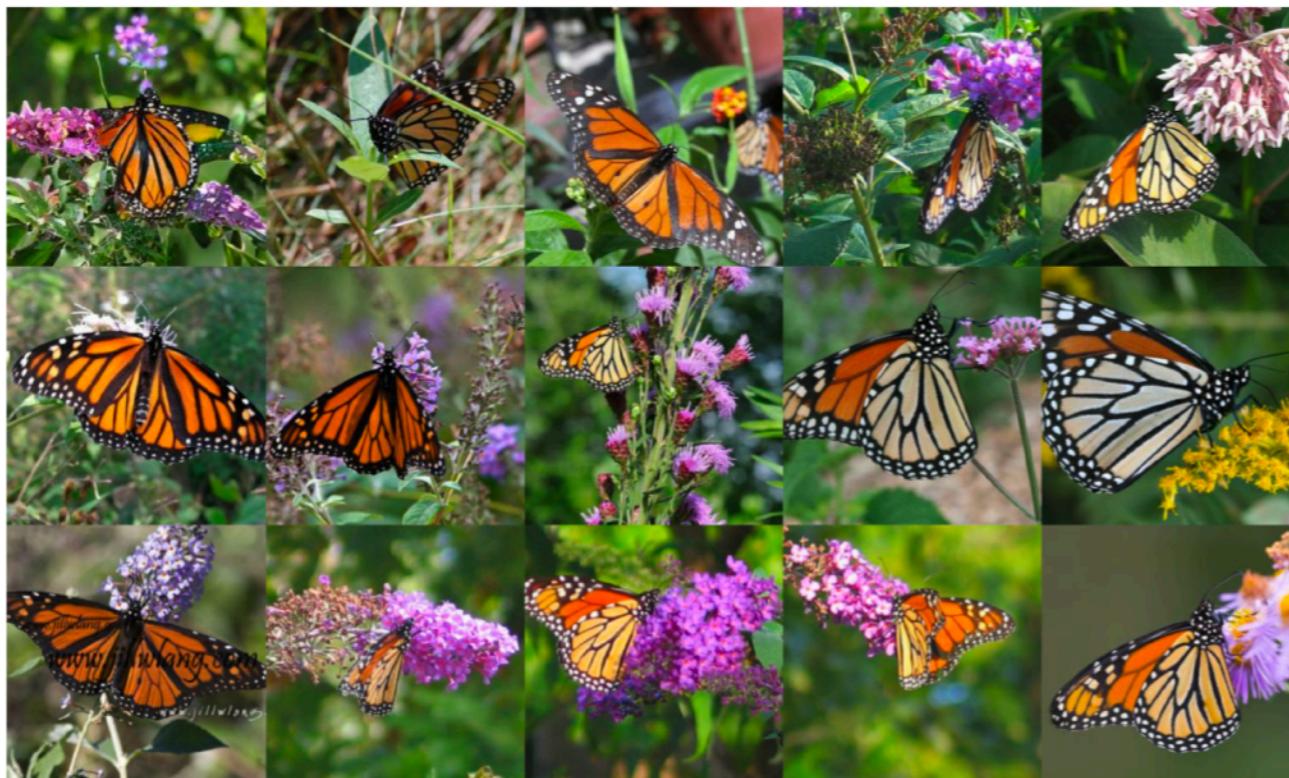
[Goodfellow *et al.*, 2014]

- Generative process
 - Draw z from p_z
 - Pass it to the generator to compute $G(z)$
- Optimization
 - Alternate between generator and discriminator
 - Very unstable process in practice

$$\begin{aligned}\min_G \max_D L(D, G) &= \mathbb{E}_{x \sim p_r(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \\ &= \mathbb{E}_{x \sim p_r(x)} [\log D(x)] + \mathbb{E}_{x \sim p_g(x)} [\log(1 - D(x))]\end{aligned}$$

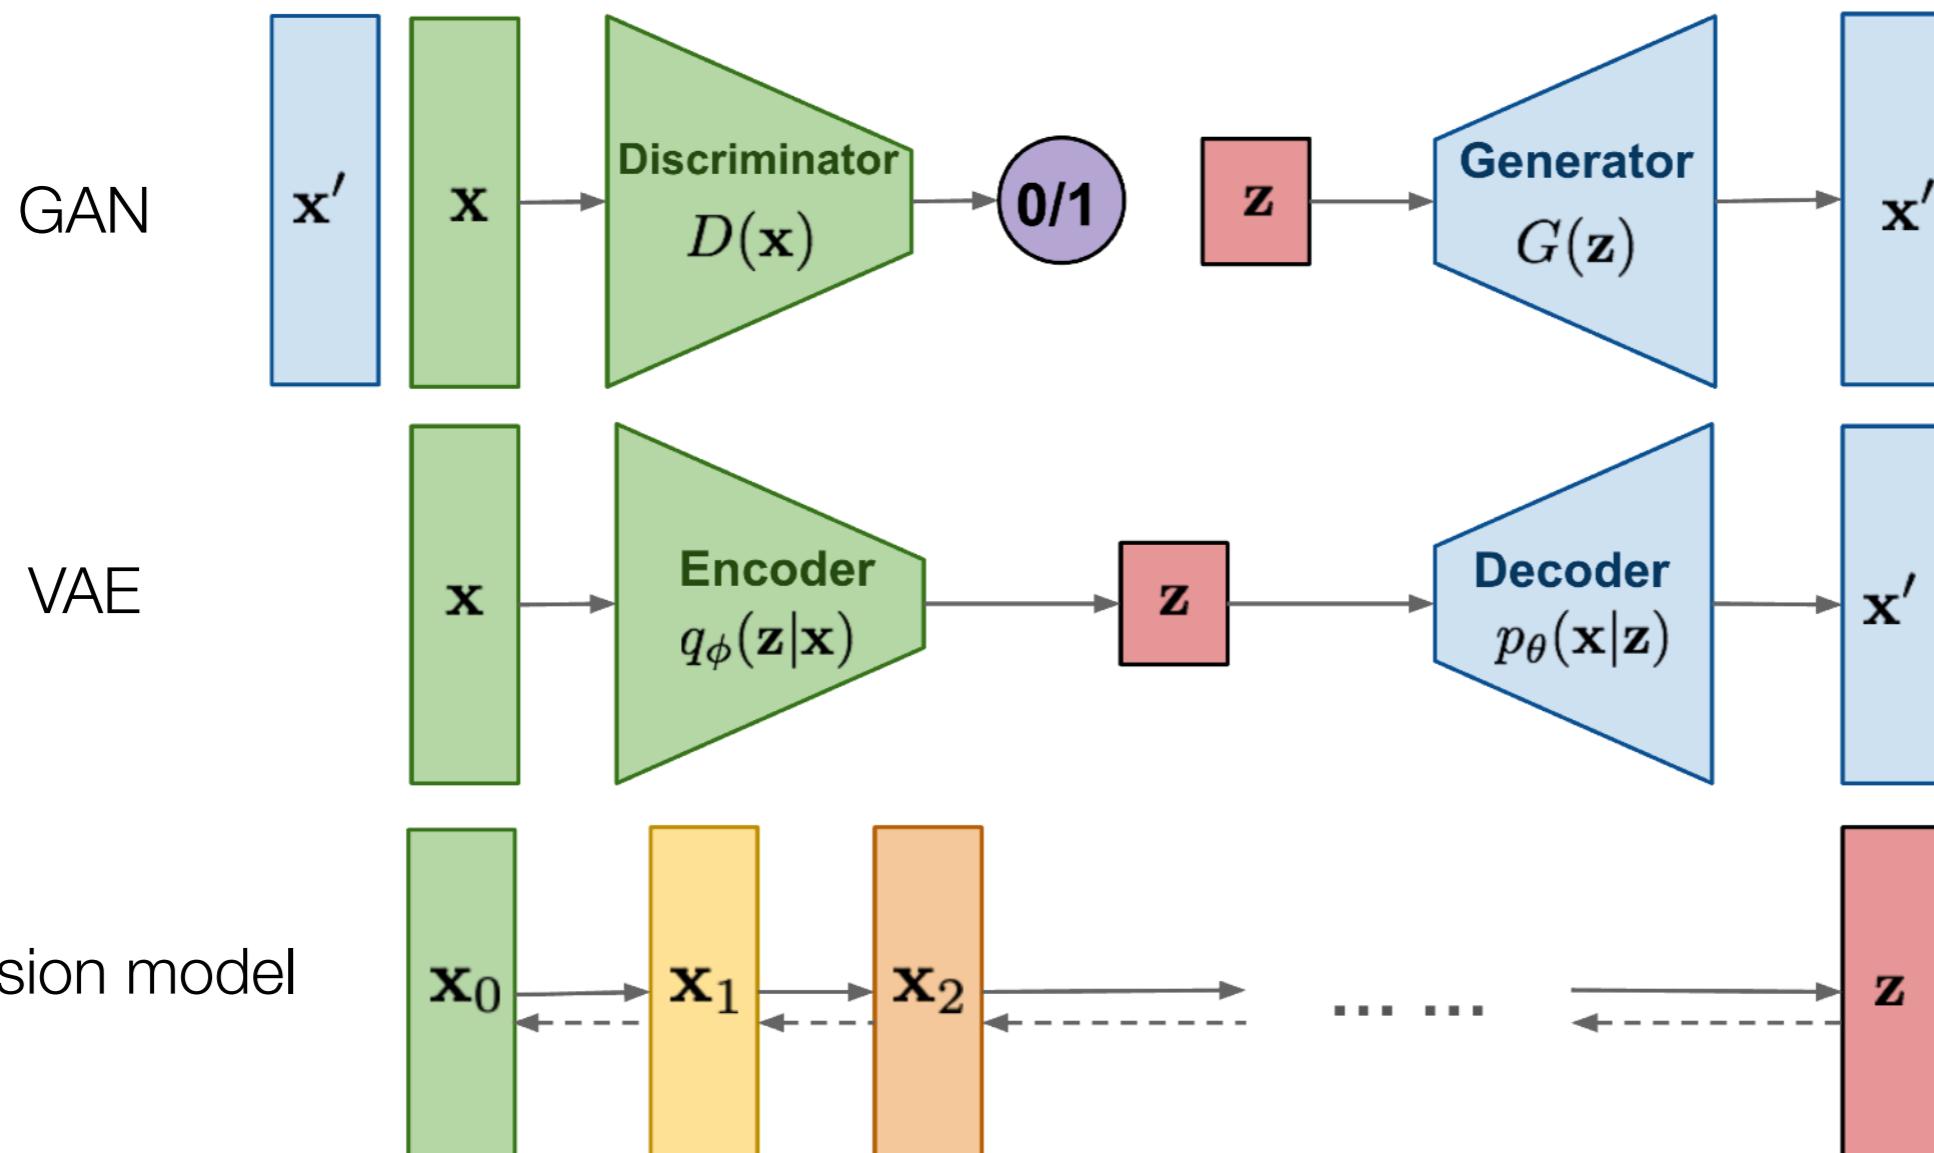
Generative Adversarial Networks

- Many variants to the original model
 - Class-conditional variants
 - Different losses
 - Different structures
- Very realistic samples generated (BigGAN, StyleGAN)



Diffusion models

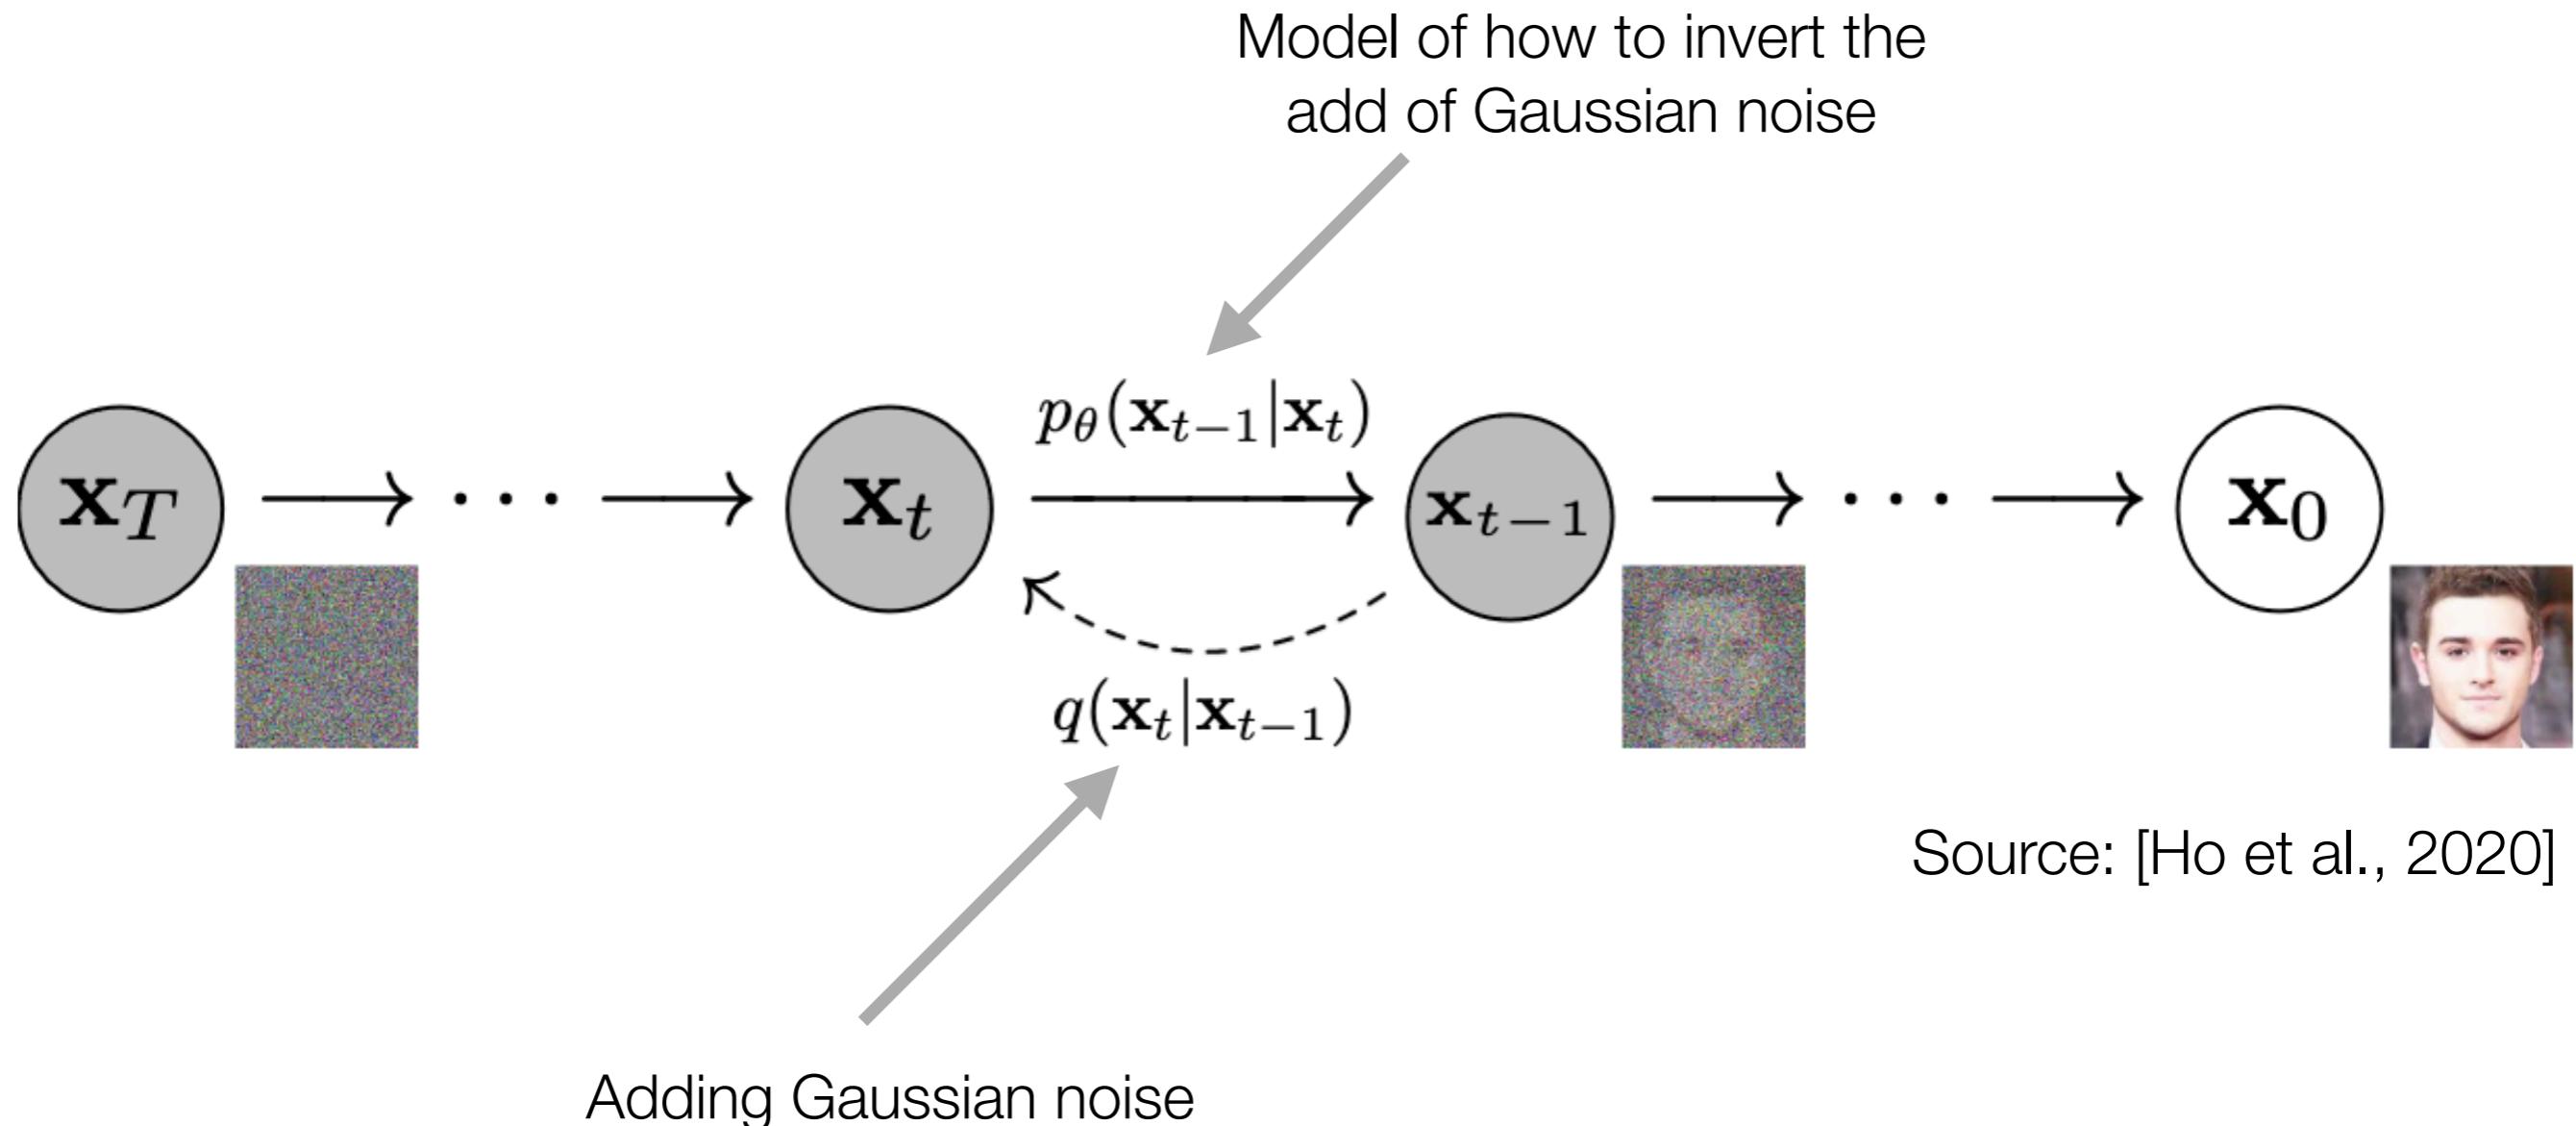
[Ho *et al.*, 2020]



Source: lilianweng.github.io

Diffusion models

[Ho *et al.*, 2020]



Diffusion models

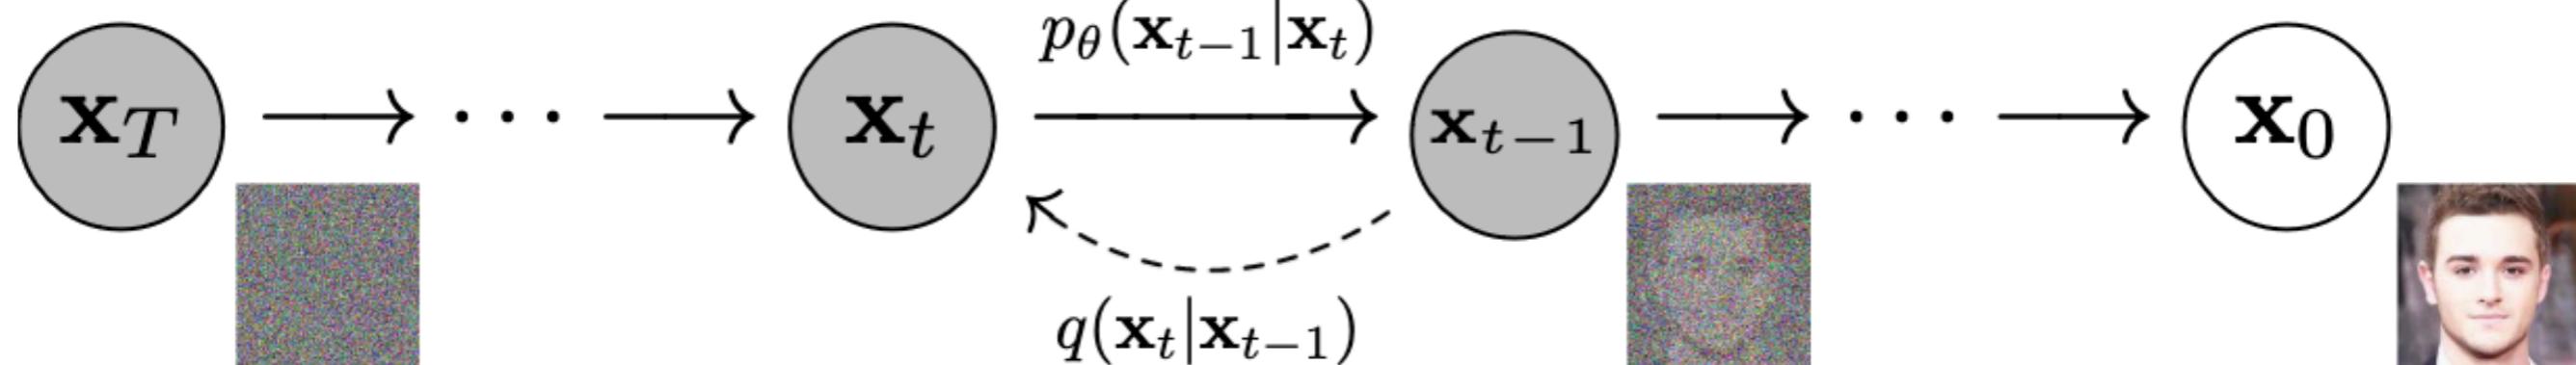
[Ho et al., 2020]

Algorithm 1 Training

```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
        $\nabla_{\theta} \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$ 
6: until converged
```

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```



Source: [Ho et al., 2020]

Conditional Flow Matching

[Lipman et al., 2023]

- Principle: learn a vector field that transports noise samples ($t=0$) to data samples ($t=1$)

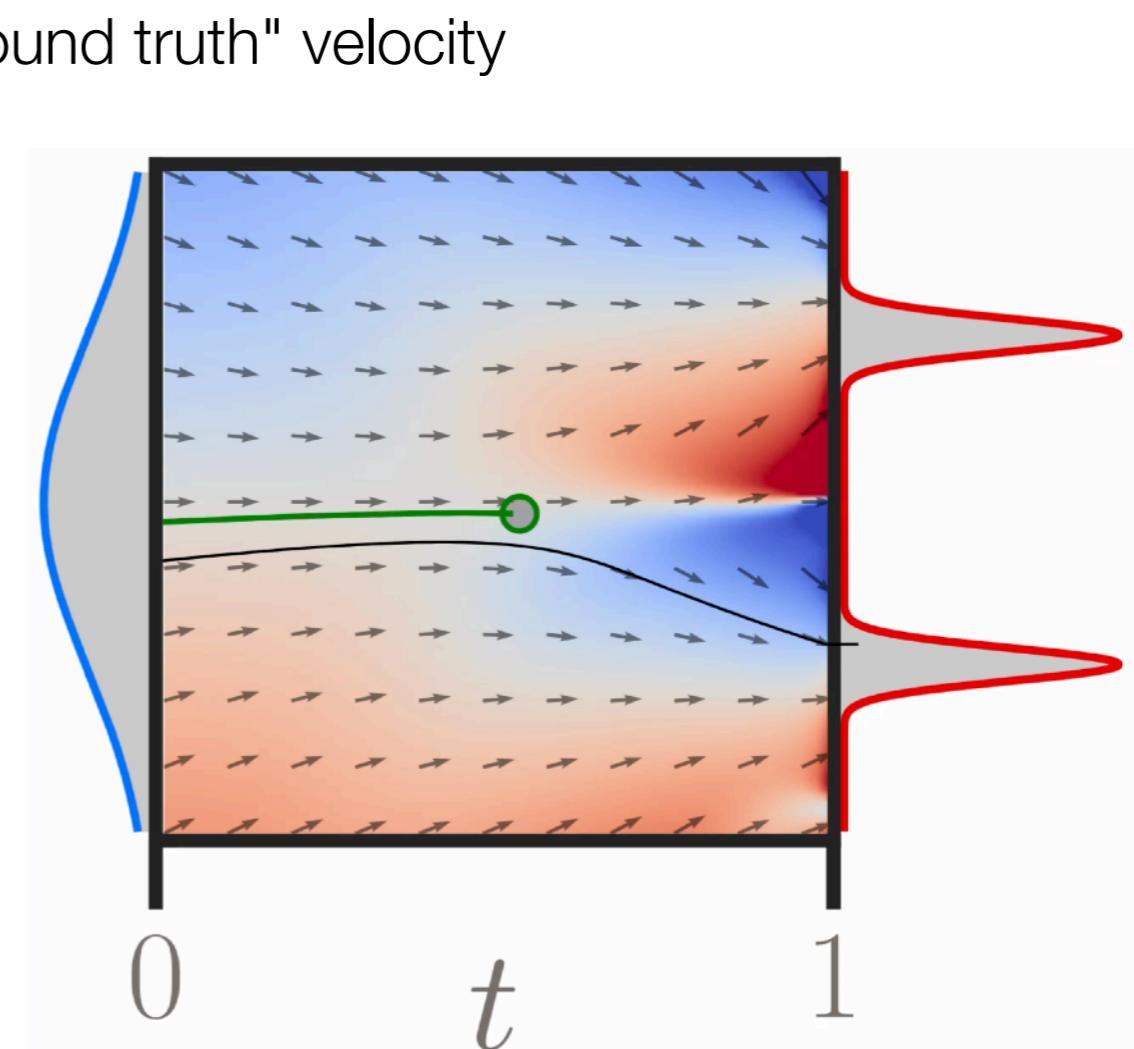
- Loss

$$\mathbb{E}_{x_0, x_1, t} \|u^\theta(x, t) - (x_1 - x_0)\|^2$$

$t \cdot x_0 + (1 - t) \cdot x_1$

- At generation time, use an Ordinary Differential Equation solver (eg. Euler scheme):

$$x_{t+\varepsilon} \leftarrow x_t + \varepsilon u^\theta(x_t, t)$$



Source: [Gagneux et al., 2025]