

Basics of keras

Romain Tavenard

Specifying model architecture

The Sequential model class

- ▶ Sequential: easiest way to code a simple model in keras
- ▶ Requirements
 - ▶ Single input (can be multidimensional)
 - ▶ Single output (can be multidimensional)
 - ▶ model is just a stack of layers (no loop in the model graph)

```
from keras.models import Sequential

model = Sequential(
    [
        # Here, we will put a list of layers
    ]
)
```

Layers used to define an MLP

- ▶ InputLayer (not mandatory, takes an input_shape argument)
- ▶ Dense (link to docs)
 - ▶ Can be used for both internal and output layers
 - ▶ activation is the activation function (str)

```
from keras.models import Sequential
from keras.layers import InputLayer, Dense

model = Sequential(
    [
        # Input layer (16-dimensional)
        InputLayer(input_shape=(16, )),
        # Hidden layer (256 neurons, ReLU activation)
        Dense(units=256, activation="relu"),
        # Output layer (1 neuron, sigmoid activation)
        Dense(units=1, activation="sigmoid")
    ]
)
```

Fitting a model

Setting optimization strategy (1/2)

- ▶ `compile` step
 - ▶ `optimizer` (required)
 - ▶ can be a string ("sgd", "adam", ...)
 - ▶ can be a keras optimizer instance (allows to set learning rate, ...)
 - ▶ `loss` (required) can be a string
 - ▶ "mse": Mean Squared Error
 - ▶ "binary_crossentropy": Cross-entropy in the binary case (single output neuron)
 - ▶ "categorical_crossentropy": Cross-entropy in the multiclass case (multiple output neurons)
 - ▶ `metrics` (optional) is a list
 - ▶ additional metrics to be reported during model training (eg. "accuracy" for classification tasks)

Setting optimization strategy (2/2)

```
model = Sequential(  
    [  
        # Layers here  
    ]  
)  
  
model.compile(  
    optimizer="sgd",  
    loss="mse",  
    metrics=[]  
)
```

Setting fit options (1/2)

- ▶ `fit` step
 - ▶ First 2 arguments: training data (required)
 - ▶ `validation_data` (optional) is a tuple
 - ▶ validation set organized as a pair (X, y)
 - ▶ `epochs`: number of epochs (required)
 - ▶ `batch_size`: number of samples per batch (strongly recommended)

Setting fit options (2/2)

```
model = Sequential(  
    [  
        # Layers here  
    ]  
)  
  
model.compile(  
    # Compile options  
)  
model.fit(X, y,  
          validation_data=(X_val, y_val),  
          epochs=10,  
          batch_size=64)
```

Callbacks

Callbacks in general

- ▶ Callbacks are ways to “interact” during training
 - ▶ eg. report statistics, early stopping, save models
 - ▶ at each minibatch or at each epoch
- ▶ A list of callbacks can be passed at fit time:

```
model.fit(X, y,  
          validation_data=(X_val, y_val),  
          epochs=10,  
          batch_size=64,  
          callbacks=[cb1, cb2, ...])
```

- ▶ See docs for usage ([link](#))
 - ▶ EarlyStopping
 - ▶ ModelCheckpoint
 - ▶ TensorBoard

History callback

- ▶ Returned by fit by default (no need to set up)
- ▶ has a history attribute
 - ▶ dict
 - ▶ losses and metrics recorded during fit at each epoch

```
h = model.fit(X, y,  
              validation_data=(X_val, y_val),  
              epochs=10,  
              batch_size=64)  
  
print(h.history)  
# {  
#   "loss": [...],  
#   "val_loss": [...],  
#   "accuracy": [...],  
#   "val_accuracy": [...]  
# }
```